

Selection Problem · Week III

Data Structures: Theory and Practice

Sunday 29th November, 2009

Johan Sejr Brinch Nielsen


Email: zerrez@diku.dk
Cpr.: 260886-2547

Dept. of Computer Science,
University of Copenhagen

1) Exercise 19.4-1

For any natural number n it is possible to generate a fibonacci heap of size n , for which its internal tree structure is linear.

- Base case $n = 0$: Insert an element. The heap is now of size 1 and linear.
- Induction step $n \rightarrow n + 1$: Suppose we have a linear fibonacci heap of size n with x as minimum. Insert three nodes a, b and c where $a < b < c$ and all three are lower than the minimum of the heap (a is the new minimum). Call extract-min to extract a . The Consolidate operation will now make a tree with b as root and x and c as children. Delete c . We now have a linear fibonacci heap of size $n + 1$.



2) Problem 19-4

A mergeable heap backed by a 2-3-4-tree. The operations could be implemented by:

- Minimum: Maintain a pointer to the smallest node to achieve $O(1)$.
- Decrease-key: Decrease the key of the leaf and update all the affected small values of its parent etc. Update the minimum-pointer if needed. This operation is $O(\log n)$.
- Insert: First, find the leaf where the node is to be inserted by choosing the appropriate child (largest small value lower than the value inserted). Now, create a new leaf for the value and add it to the parent of the found leaf. If the parent is full do splitting as in 2-3-4-trees. Update the small values and the minimum-pointer if the inserted value is a new minimum. Time: $O(\log n)$
- Delete: This is done as in an 2-3-4-tree. First delete the node. If its parent becomes empty perform the appropriate merging. Update the appropriate small values. If the deleted key was the minimum, scan the tree for the new minimum. This operation is $O(\log n)$ time.

- **Extract-min:** Call `Delete` on the minimum (given by the pointer). This is $O(\log n)$.
- **Union:** How to make a new heap H from two existing heaps H_0 and H_1 . If the two H_0 and H_1 have equal heights, make a new root that these as children. Set its small value to the minimum of the two heaps. This case takes $O(1)$.

If H_0 is tallest, find some node in H_0 whose subtrees are the same height as H_1 and add H_1 as a child.

Finding H_1 's new parent in H_0 is $O(\log(\max(n, m))) = O(\log(n + m))$.

Stubbed C++ code

File: `heap.h++`

```
using namespace std;

template <class T>
class HeapNode {
public:
    // only internal nodes:
    T small;
    HeapNode<T> *children[3];

    // only leafs:
    int key_count;
    T keys[3];
};

template <class T>
class Heap {
public:
    HeapNode<T>* search(T key);
    void insert(T key);
    void delete_key(HeapNode<T> *node);
    HeapNode<T>* minimum();
    void decrease_key(HeapNode<T> *node, T new_key);
    HeapNode<T>* extract_min();
    Heap<T>* merge(Heap<T> *heap);
    void print(HeapNode<T> *node);
private:
    HeapNode<T>* root;
    HeapNode<T>* minimum;
};
```

File: `heap.c++`

```
#include "heap.h++"
#include <stdlib.h>
```

```

#include<stdio.h>

using namespace std;

template <class T>
void Heap<T>::insert(T key) {}

template <class T>
HeapNode<T>* Heap<T>::search(T key) {}

template <class T>
void delete_key(HeapNode<T> *node) {}

template <class T>
HeapNode<T>* minimum() {}

template <class T>
HeapNode<T>* extract_min() {}

template <class T>
Heap<T>* merge(Heap<T> *heap) {}

template <class T>
void print(HeapNode<T> *node) {}

```

3) Binomial Heaps

I have investigated some running times in binomial heaps:

- n inserts: $O(n)$
- m deletes: $O(m \log m)$
- n inserts then m deletes: $O(n + m \log m)$

I convinced that the running time of a mixed sequence of inserts and deletes are worst case $O(N \log N)$ where N is the total number of operations.



4) Pennants

A pennant is a heap ordered binary tree with a size of 2^r where r is the rank. The root node has no left child and its right subtree is a perfect binary tree. The minimum value of the tree is stored in the root node. Consider now an operation that merges to trees of same rank:

Merge(s , t)

Two pennants can be merged by:

1. Initialize a new pennant p .
2. If s is the tree with the smallest root and t the other tree then extract the root of s and let it be the new root of p .

Figure 0.1: Example of two pennants s and t Figure 0.2: The two pennants s and t merged

3. If now $\text{Root}(t) < \text{Root}(s)$ then set $\text{Right}(\text{Root}(p)) = t$ and $\text{Left}(t) = s$. p is now the merged tree. This case requires $O(1)$ time.

If instead $\text{Root}(t) > \text{Root}(s)$ then swap the two and set $\text{Right}(\text{Root}(p)) = t$. Set $\text{Left}(t) = s$. The problem now is that $\text{Left}(t)$ may now be a heap-ordered tree. It's root could be larger than its children. We therefore need to push the down until it reaches a smaller node or the end. This case requires $O(r)$ time (This is the case illustrated in the example).

The worst case running time of the Merge operation is thereby $O(r)$ where r is the rank of the two trees (each of size 2^r).

The worst case situation when performing a Union operation on two pennant queues of size S is when both queues has a pennant for each rank r . In this case S Merge operations are performed. The worst case running time of this operation is thereby:

$$O\left(\sum_{r=0}^S r\right) = O\left(\frac{S^2 + S}{2}\right) = O(S^2) = O(\log^2(N))$$

Where N is the number of elements in the united queues.

5) Fibonacci Heaps

A normal fibonacci heap has the following worst case running times for the priority queue operations:

- Insert: $O(1)$
- Find-min: $O(1)$
- Extract-min: $O(n)$

The worst case of Extract-min is bounded by the maximum of children the minimum-node has. If we give an upper bound of $O(\log n)$ on this number the operation would run in worst case $O(\log n)$. My idea is to move some of the work from the Extract-min onto the Insert operation which currently runs in $O(1)$. It goes as follows:

- Set an upper bound of $2 \log n$ on the children any node can have.
- Keep track of the count of a node's children when inserting and extracting using a counter.
- If a node has $2 \log n$ children when inserting a new child, cut the $\log n$ of them and make them roots. Run consolidate.

If one could actually maintain an invariant that guaranteed that no node has more than $O(\log n)$ children, then Extract-min would have $O(\log n)$ worst case.